

In [1]: `import numpy as np`

READING THE DATA

In [7]: `import pandas as pd
df = pd.read_csv("restaurant_data.csv")
df`

Out[7]:

	Name	Location	Cuisine	Rating	Seating Capacity	Average Meal Price	Marketing Budget	Social Media Followers	Chef Experience Years	Number of Reviews	Avg Review Length	Ambience Score	Service Quality Score	Parking Availability	Weekend Reservations	Weekday Reservations	Revenue
0	Restaurant 0	Rural	Japanese	4.0	38	73.98	2224	23406	13	185	161.924906	1.3	7.0	Yes	13	4	638945.52
1	Restaurant 1	Downtown	Mexican	3.2	76	28.11	4416	42741	8	533	148.759717	2.6	3.4	Yes	48	6	490207.83
2	Restaurant 2	Rural	Italian	4.7	48	48.29	2796	37285	18	853	56.849189	5.3	6.7	No	27	14	541368.62
3	Restaurant 3	Rural	Italian	4.4	34	51.55	1167	15214	13	82	205.433265	4.6	2.8	Yes	9	17	404556.80
4	Restaurant 4	Downtown	Japanese	4.9	88	75.98	3639	40171	9	78	241.681584	8.6	2.1	No	37	26	1491046.35
...
8363	Restaurant 8363	Suburban	Indian	3.4	54	34.85	1102	11298	11	380	253.919515	9.5	5.0	Yes	37	0	434653.45
8364	Restaurant 8364	Rural	Indian	3.7	49	36.88	1988	20432	9	713	175.590194	2.7	2.6	No	37	21	414977.92
8365	Restaurant 8365	Downtown	Italian	4.7	88	46.87	5949	63945	6	436	222.953647	4.8	1.7	Yes	83	21	930395.87
8366	Restaurant 8366	Rural	American	3.1	31	44.53	707	7170	1	729	178.482851	6.1	2.1	No	6	21	311493.48
8367	Restaurant 8367	Rural	Japanese	4.0	33	71.07	2003	24268	8	197	151.838065	5.9	7.5	Yes	5	12	534142.98

8368 rows × 17 columns

In [8]: `df.info()`



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8368 entries, 0 to 8367
Data columns (total 17 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Name              8368 non-null    object  
 1   Location          8368 non-null    object  
 2   Cuisine            8368 non-null    object  
 3   Rating             8368 non-null    float64 
 4   Seating Capacity  8368 non-null    int64  
 5   Average Meal Price 8368 non-null    float64 
 6   Marketing Budget  8368 non-null    int64  
 7   Social Media Followers 8368 non-null    int64  
 8   Chef Experience Years 8368 non-null    int64  
 9   Number of Reviews  8368 non-null    int64  
 10  Avg Review Length 8368 non-null    float64 
 11  Ambience Score   8368 non-null    float64 
 12  Service Quality Score 8368 non-null    float64 
 13  Parking Availability 8368 non-null    object  
 14  Weekend Reservations 8368 non-null    int64  
 15  Weekday Reservations 8368 non-null    int64  
 16  Revenue            8368 non-null    float64 
dtypes: float64(6), int64(7), object(4)
memory usage: 1.1+ MB
```

In [10]: df.shape

Out[10]: (8368, 17)

Checking if any row or column has null value or not


In [11]: df.isna().sum()

```
Name          0
Location      0
Cuisine        0
Rating         0
Seating Capacity 0
Average Meal Price 0
Marketing Budget 0
Social Media Followers 0
Chef Experience Years 0
Number of Reviews 0
Avg Review Length 0
Ambience Score 0
Service Quality Score 0
Parking Availability 0
Weekend Reservations 0
Weekday Reservations 0
Revenue         0
dtype: int64
```


Making charts and plots using matplotlib and seaborn libs

```
In [22]: import matplotlib.pyplot as plt  
import seaborn as sns
```


```
plt.figure(figsize=(10,6))  
sns.violinplot(data=df, x = 'Chef Experience Years', y = 'Revenue')  
plt.title('Chef Experience Years vs Revenue')  
plt.show()
```



```
In [21]: plt.figure(figsize=(10,6))  
plt.title('Location vs Revenue')  
sns.boxplot(data=df, x = 'Location', y = 'Revenue')  
plt.show()
```



```
In [23]: plt.figure(figsize=(10,6))
plt.title('Rating vs Revenue')
sns.scatterplot(data=df, x = 'Rating', y = 'Revenue')
plt.show()
```




```
In [26]: plt.figure(figsize=(10,6))
sns.lmplot(x='Average Meal Price', y='Revenue', data=df)
plt.title('Average Meal Price vs Revenue')
plt.show()
```

<Figure size 1000x600 with 0 Axes>




```
In [27]: plt.figure(figsize=(10,6))
sns.histplot(df['Marketing Budget'], bins=30)
plt.title('Marketing Budget Distribution')
plt.xlabel('Marketing Budget')
plt.ylabel('Frequency')
plt.show()
```

Marketing Budget Distribution




```
In [32]: plt.figure(figsize=(10,6))
sns.scatterplot(x='Marketing Budget', y='Revenue', data=df)
plt.title('Marketing Budget vs Revenue')
plt.xlabel('Marketing Budget')
plt.ylabel('Revenue')
plt.show()
```



```
In [42]: plt.figure(figsize=(10,6))
sns.jointplot(x='Weekend Reservations', y='Revenue', data=df, kind='kde')
plt.title('Weekend Reservations vs Revenue')
plt.xlabel('Weekend Reservation')
plt.ylabel('Revenue')
plt.show()
```

<Figure size 1000x600 with 0 Axes>



```
In [47]: string_columns = df.select_dtypes(include=['object'])
string_columns.head()
```

```
Out[47]:
```

	Name	Location	Cuisine	Parking Availability
0	Restaurant 0	Rural	Japanese	Yes
1	Restaurant 1	Downtown	Mexican	Yes
2	Restaurant 2	Rural	Italian	No
3	Restaurant 3	Rural	Italian	Yes
4	Restaurant 4	Downtown	Japanese	No



READING THE DATA/ TAKING INSIGHTS

```
In [55]: df['Parking Availability'].unique()
```

```
Out[55]: array(['Yes', 'No'], dtype=object)
```

```
In [56]: df['Parking Availability'] = df['Parking Availability'].replace({'Yes': 1, 'No': 0})
df['Parking Availability'].head()
```

```
Out[56]: 0    1
1    1
2    0
3    1
4    0
Name: Parking Availability, dtype: int64
```

```
In [57]: df['Location'].unique()
```

```
Out[57]: array(['Rural', 'Downtown', 'Suburban'], dtype=object)
```

```
In [58]: df['Location'].value_counts()
```

```
Out[58]: Downtown    2821
Suburban     2785
Rural        2762
Name: Location, dtype: int64
```

```
In [59]: location_variables = pd.get_dummies(df['Location'], dtype=int)
location_variables.head()
```

```
Out[59]:   Downtown  Rural  Suburban
0          0      1        0
1          1      0        0
2          0      1        0
3          0      1        0
4          1      0        0
```

```
In [60]: updated_df = pd.concat([df, location_variables], axis=1)
updated_df.head()
```

Restaurant Revenue Prediction

Out[60]:

	Name	Location	Cuisine	Rating	Seating Capacity	Average Meal Price	Marketing Budget	Social Media Followers	Chef Experience Years	Number of Reviews	Avg Review Length	Ambience Score	Service Quality Score	Parking Availability	Weekend Reservations	Weekday Reservations	Revenue	Dov
0	Restaurant 0	Rural	Japanese	4.0	38	73.98	2224	23406	13	185	161.924906	1.3	7.0	1	13	4	638945.52	
1	Restaurant 1	Downtown	Mexican	3.2	76	28.11	4416	42741	8	533	148.759717	2.6	3.4	1	48	6	490207.83	
2	Restaurant 2	Rural	Italian	4.7	48	48.29	2796	37285	18	853	56.849189	5.3	6.7	0	27	14	541368.62	
3	Restaurant 3	Rural	Italian	4.4	34	51.55	1167	15214	13	82	205.433265	4.6	2.8	1	9	17	404556.80	
4	Restaurant 4	Downtown	Japanese	4.9	88	75.98	3639	40171	9	78	241.681584	8.6	2.1	0	37	26	1491046.35	

In [61]: `updated_df.shape`

Out[61]: `(8368, 20)`

In [62]: `print(updated_df['Downtown'].sum())
print(updated_df['Rural'].sum())
print(updated_df['Suburban'].sum())`

2821
2762
2785

In [63]: `updated_df = updated_df.drop('Location', axis=1)
updated_df.head()`

Out[63]:

	Name	Cuisine	Rating	Seating Capacity	Average Meal Price	Marketing Budget	Social Media Followers	Chef Experience Years	Number of Reviews	Avg Review Length	Ambience Score	Service Quality Score	Parking Availability	Weekend Reservations	Weekday Reservations	Revenue	Downtown	Ru
0	Restaurant 0	Japanese	4.0	38	73.98	2224	23406	13	185	161.924906	1.3	7.0	1	13	4	638945.52	0	
1	Restaurant 1	Mexican	3.2	76	28.11	4416	42741	8	533	148.759717	2.6	3.4	1	48	6	490207.83	1	
2	Restaurant 2	Italian	4.7	48	48.29	2796	37285	18	853	56.849189	5.3	6.7	0	27	14	541368.62	0	
3	Restaurant 3	Italian	4.4	34	51.55	1167	15214	13	82	205.433265	4.6	2.8	1	9	17	404556.80	0	
4	Restaurant 4	Japanese	4.9	88	75.98	3639	40171	9	78	241.681584	8.6	2.1	0	37	26	1491046.35	1	

```
In [64]: updated_df.shape
```

```
Out[64]: (8368, 19)
```

```
In [65]: string_columns = updated_df.select_dtypes(include=['object'])
string_columns.head()
```

```
Out[65]:
```

	Name	Cuisine
0	Restaurant 0	Japanese
1	Restaurant 1	Mexican
2	Restaurant 2	Italian
3	Restaurant 3	Italian
4	Restaurant 4	Japanese

```
In [85]: df['Cuisine'].unique()
```

```
Out[85]: array(['Japanese', 'Mexican', 'Italian', 'Indian', 'French', 'American'],
      dtype=object)
```

```
In [86]: df['Cuisine'].value_counts()
```

```
Out[86]:
```

French	1433
American	1416
Italian	1413
Mexican	1393
Indian	1369
Japanese	1344

Name: Cuisine, dtype: int64

```
In [66]: numeric_cuisine_values = pd.get_dummies(df['Cuisine'], dtype=int)
numeric_cuisine_values.head()
```

```
Out[66]:
```

	American	French	Indian	Italian	Japanese	Mexican
0	0	0	0	0	1	0
1	0	0	0	0	0	1
2	0	0	0	1	0	0
3	0	0	0	1	0	0
4	0	0	0	0	1	0

```
In [67]: print("counts for French: ", numeric_cuisine_values['French'].sum())
print("counts for American: ", numeric_cuisine_values['American'].sum())
print("counts for Italian: ", numeric_cuisine_values['Italian'].sum())
print("counts for Mexican: ", numeric_cuisine_values['Mexican'].sum())
print("counts for Indian: ", numeric_cuisine_values['Indian'].sum())
print("counts for Japanese: ", numeric_cuisine_values['Japanese'].sum())
```

```
counts for French: 1433
counts for American: 1416
counts for Italian: 1413
counts for Mexican: 1393
counts for Indian: 1369
counts for Japanese: 1344
```

In [68]: `updated_df.shape`

Out[68]: (8368, 19)

In [69]: `updated_df = pd.concat([updated_df, numeric_cuisine_values], axis=1)`
`updated_df.head()`

Out[69]:

	Name	Cuisine	Rating	Seating Capacity	Average Meal Price	Marketing Budget	Social Media Followers	Chef Experience Years	Number of Reviews	Avg Review Length	...	Revenue	Downtown	Rural	Suburban	American	French	Indian	Italian	Japanese
0	Restaurant 0	Japanese	4.0	38	73.98	2224	23406	13	185	161.924906	...	638945.52	0	1	0	0	0	0	0	0
1	Restaurant 1	Mexican	3.2	76	28.11	4416	42741	8	533	148.759717	...	490207.83	1	0	0	0	0	0	0	0
2	Restaurant 2	Italian	4.7	48	48.29	2796	37285	18	853	56.849189	...	541368.62	0	1	0	0	0	0	0	1
3	Restaurant 3	Italian	4.4	34	51.55	1167	15214	13	82	205.433265	...	404556.80	0	1	0	0	0	0	0	1
4	Restaurant 4	Japanese	4.9	88	75.98	3639	40171	9	78	241.681584	...	1491046.35	1	0	0	0	0	0	0	0

5 rows × 25 columns

In [89]: `updated_df = updated_df.drop('Cuisine', axis=1)`
`updated_df.head()`

Restaurant Revenue Prediction

Out[89]:

	Rating	Seating Capacity	Average Meal Price	Marketing Budget	Social Media Followers	Chef Experience Years	Number of Reviews	Avg Review Length	Ambience Score	Service Quality Score	...	Revenue	Downtown	Rural	Suburban	American	French	Indian	Italian	Japanese
0	4.0	38	73.98	2224	23406	13	185	161.924906	1.3	7.0	...	638945.52	0	1	0	0	0	0	0	
1	3.2	76	28.11	4416	42741	8	533	148.759717	2.6	3.4	...	490207.83	1	0	0	0	0	0	0	
2	4.7	48	48.29	2796	37285	18	853	56.849189	5.3	6.7	...	541368.62	0	1	0	0	0	0	1	
3	4.4	34	51.55	1167	15214	13	82	205.433265	4.6	2.8	...	404556.80	0	1	0	0	0	0	1	
4	4.9	88	75.98	3639	40171	9	78	241.681584	8.6	2.1	...	1491046.35	1	0	0	0	0	0	0	

5 rows × 23 columns

In [90]: `updated_df.head()`

Out[90]:

	Rating	Seating Capacity	Average Meal Price	Marketing Budget	Social Media Followers	Chef Experience Years	Number of Reviews	Avg Review Length	Ambience Score	Service Quality Score	...	Revenue	Downtown	Rural	Suburban	American	French	Indian	Italian	Japanese
0	4.0	38	73.98	2224	23406	13	185	161.924906	1.3	7.0	...	638945.52	0	1	0	0	0	0	0	
1	3.2	76	28.11	4416	42741	8	533	148.759717	2.6	3.4	...	490207.83	1	0	0	0	0	0	0	
2	4.7	48	48.29	2796	37285	18	853	56.849189	5.3	6.7	...	541368.62	0	1	0	0	0	0	1	
3	4.4	34	51.55	1167	15214	13	82	205.433265	4.6	2.8	...	404556.80	0	1	0	0	0	0	1	
4	4.9	88	75.98	3639	40171	9	78	241.681584	8.6	2.1	...	1491046.35	1	0	0	0	0	0	0	

5 rows × 23 columns

In [91]: `data = updated_df
data.shape`Out[91]: `(8368, 23)`In [92]: `np.random.seed(42)`In [93]: `x = data.drop('Revenue', axis=1)
x.head(2)`

Restaurant Revenue Prediction

Out[93]:	Rating	Seating Capacity	Average Meal Price	Marketing Budget	Social Media Followers	Chef Experience Years	Number of Reviews	Avg Review Length	Ambience Score	Service Quality Score	...	Weekday Reservations	Downtown	Rural	Suburban	American	French	Indian	Italian	Japanese
0	4.0	38	73.98	2224	23406	13	185	161.924906	1.3	7.0	...	4	0	1	0	0	0	0	0	
1	3.2	76	28.11	4416	42741	8	533	148.759717	2.6	3.4	...	6	1	0	0	0	0	0	0	

2 rows × 22 columns

In [94]:

```
y = data['Revenue']
y.head()
```

Out[94]:

```
0    638945.52
1    490207.83
2    541368.62
3    404556.80
4    1491046.35
Name: Revenue, dtype: float64
```

Let's train the model

Starting with Linear Regression Model

In [101...]

```
from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.4)

print(f"x_train:{x_train.shape}, x_test:{x_test.shape} y_train:{y_train.shape}, y_test:{y_test.shape}")
x_train:(5020, 22), x_test:(3348, 22) y_train:(5020,), y_test:(3348,)
```

In [102...]

```
from sklearn.linear_model import LinearRegression

model = LinearRegression()
model.fit(x_train, y_train)
```

Out[102]:

```
LinearRegression()
LinearRegression()
```

In [103...]

```
y_pred = model.predict(x_test)
```


In [104...]

```
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
mse = mean_squared_error(y_test, y_pred)
r2Score = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
print("Mean Squared Error: ", mse)
print("R squared Error: ", r2Score)
print("Mean Absolute Error: ", mae)
```

```
print("R2 score for training: ", model.score(x_train,y_train))
print(model.score(x_test, y_test))
```

Mean Squared Error: 3033168683.267786
R squared Error: 0.9572660348922483
Mean Absolute Error: 41909.96555274624
R2 score for training: 0.9584679045757779
0.9572660348922483

```
In [114...]: plt.scatter(y_pred, y_test, color='yellow')
plt.xlabel("predicted")
plt.ylabel("Actual")
plt.title("Actual vs Predicted values")
plt.show()
```




```
In [118...]: from sklearn.decomposition import PCA
pca = PCA(n_components=3)
x_train_pca = pca.fit_transform(x_train)
x_test_pca = pca.transform(x_test)

fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111, projection='3d')

# actual values
ax.scatter(x_test_pca[:, 0], x_test_pca[:, 1], y_test, color='violet', label='Actual')
```

```
# predicted values
ax.scatter(x_test_pca[:, 0], x_test_pca[:, 1], y_pred, color='red', label='Predicted')
ax.set_xlabel('PCA Component 1')
ax.set_ylabel('PCA Component 2')
ax.set_zlabel('Y')
ax.set_title('Linear Regression: Actual vs Predicted (3D Plot)')
ax.legend()
plt.show()
```

Linear Regression: Actual vs Predicted (3D Plot)



Linear Regression Accuracy = 95.72%

Do these for Random Forest also !

```
In [121...]  
from sklearn.ensemble import RandomForestRegressor  
clf = RandomForestRegressor()  
clf.fit(x_train, y_train)
```

```
Out[121]:  
RandomForestRegressor()  
RandomForestRegressor()
```


```
In [109...]  
y2_pred = clf.predict(x_test)
```

```
In [110...]  
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error  
mse = mean_squared_error(y_test, y2_pred)  
r2Score = r2_score(y_test, y2_pred)  
mae = mean_absolute_error(y_test, y2_pred)  
print("Mean Squared Error: ", mse)  
print("R squared Error: ", r2Score)  
print("Mean Absolute Error: ", mae)  
print("R2 score for training: ", clf.score(x_train,y_train))  
print("R2 score for testing: ", clf.score(x_test, y_test))  
clf.score(x_train, y_train)
```

Mean Squared Error: 66913733.25353832
R squared Error: 0.9990572600996904
Mean Absolute Error: 6361.774014516131
R2 score for training: 0.9998523578720899
R2 score for testing: 0.9990572600996904
0.9998523578720899


```
Out[110]:
```

```
In [119...]  
plt.scatter(y_test, y2_pred , color='yellow')  
plt.xlabel("Actual")  
plt.ylabel("Predicted")  
plt.title("Actual vs Predicted")  
plt.show()
```



```
In [120...]  
from sklearn.decomposition import PCA  
pca = PCA(n_components=3)  
x_train_pca = pca.fit_transform(x_train)  
x_test_pca = pca.transform(x_test)  
  
fig = plt.figure(figsize=(10, 8))  
ax = fig.add_subplot(111, projection='3d')  
  
# actual values  
ax.scatter(x_test_pca[:, 0], x_test_pca[:, 1], y_test, color='violet', label='Actual')  
  
# predicted values  
ax.scatter(x_test_pca[:, 0], x_test_pca[:, 1], y2_pred, color='red', label='Predicted')  
  
ax.set_xlabel('PCA Component 1')  
ax.set_ylabel('PCA Component 2')  
ax.set_zlabel('Y')  
ax.set_title('Linear Regression: Actual vs Predicted (3D Plot)')  
ax.legend()  
  
plt.show()
```

Linear Regression: Actual vs Predicted (3D Plot)



```
In [113]:  
from sklearn.ensemble import RandomForestRegressor  
clf = RandomForestRegressor(n_estimators=120)  
clf.fit(x_train, y_train)  
clf.score(x_train, y_train)
```

```
Out[113]: 0.9998547400585148
```

Random Forest Accuracy = 99.98%

conclusion

Accuracy of the selected two models are-

1. Linear Regression Model: 95.72%
2. Random Forest Model: 99.98%

So selecting 'Random Forest Model' as a model deployment.

This marks as the end of the project.

----- Thank You -----

In []:

